Spatiotemporal Modeling of Extreme Events

Samuel Morris (NC State), Brian Reich (NC State), and Dan Cooley (Colorado State)

JSM 2014, Boston

Motivation

- Average behavior is important to understand, but it does not paint the whole picture.
  - e.g. When constructing river levees, engineers need to be able to estimate a 100-year or 1000-year flood levels.
- Spatial extreme borrows strength across space to estimate return levels and make predictions at unknown locations..
- In geostatistical analysis, kriging uses spatial correlation for prediction.
- Want to explore similar methods for extremes.

Standard analysis - Block maxima

- Uses yearly maxima
- Discards many observations
- Models are fit using the generalized extreme value distribution
- For a spatial analysis, max-stable processes give an appropriate limiting distribution

Standard analysis - Peaks over threshold

- Incorporates more data than block maxima
- Select a threshold, T, and fit data above the threshold using the generalized Pareto distribution
- Temporal dependence may be an issue between observations (e.g. flood levels don't dissipate overnight)
- Generalized Pareto distribution (GPD) is used for the exceedances.

Multivariate analysis

- Multivariate max-stable and GPD models have nice features, but they are
  - computationally challenging to work with
  - joint distribution only available in low dimension
- Pairwise likelihood approach (Huser and Davison, 2014)

Model objectives

- Our objective is to build a model that
  - has marginal distribution with a flexible tail
  - has asymptotic spatial dependence
  - has computation on the order of Gaussian models for large space-time datasets

Thresholding data

- We threshold the observed data at a high threshold T.
- Thresholded data:

$$Y_t^*(\mathbf{s}) = \begin{cases} Y_t(\mathbf{s}) & Y_t(\mathbf{s}) > T \\ T & Y_t(\mathbf{s}) \leq T \end{cases}$$

- Allows tails of the distribution to speak for themselves.

χ coefficient

- The χ coefficient is a measure of extremal dependence
- Specifically, we focus on $\chi(\mathbf{h})$ for the upper tail given by

$$\chi(\mathbf{h}) = \lim_{c \to \infty} \Pr(Y(\mathbf{s}) > c \mid Y(\mathbf{s} + \mathbf{h}) > c)$$

- If $\chi(\mathbf{h}) = 0$, then observations are asymptotically independent at distance $\mathbf{h}$.
- We expect $\lim_{\mathbf{h} \to \infty} \chi(h) = 0$.

Gaussian spatial model

- In geostatistics $Y(\mathbf{s})$ are often modeled using a Gaussian process with mean function $\mu(\mathbf{s})$ and covariance function $\rho(\mathbf{h})$.
- Model properties:
  - Nice computing properties (closed-form likelihood)
  - For a Gaussian spatial model $\lim_{c\to\infty}\chi(c) = 0$ regardless of the strength of the correlation in the bulk of the distribution.
  - Tail is not flexible (Gaussian is light tailed)

Spatial skew-t distribution

- Assume observed data $Y_t(\mathbf{s})$ come from a skew-t (Zhang and El-Shaarawi, 2012)

$$Y_t(\mathbf{s}) = X_t(\mathbf{s})\beta + \alpha z_t + v_t(\mathbf{s})$$

  where

  - $\alpha \in \mathcal{R}$ controls the skewness
  - $z_t \overset{iid}{\sim} N_{(0,\infty)}(0, \sigma_t^2)$ is a random effect
  - $v_t(\mathbf{s})$ is a Gaussian process with variance σ_t^2 and Matérn correlation
  - $\sigma_t^2 \overset{iid}{\sim} \text{IG}(a, b)$

Spatial skew-t distribution

- Conditioned on z_t and σ_t^2, $Y_t(\mathbf{s})$ is a Gaussian spatial model
- Can use standard geostatistical methods to fit this model.
- Predictions can be made through Kriging.
- Marginalizing over z_t and σ_t^2 (via MCMC),

$$Y_t(\mathbf{s}) \sim \text{skew-t}(\mu, \Sigma^*, \alpha, \text{df} = 2a)$$

where

- μ is the location
- a, b are the IG parameters for σ_t^2
- $\Sigma^* = \frac{b}{a}\Sigma$ is a scale matrix, and Σ is a Matérn covariance matrix
- $\alpha \in \mathcal{R}$ controls the skewness

Spatial skew-t distribution

- Model properties
  - Has flexible tail controlled by skewness α and degrees of freedom $2a$
  - For a skew-t distribution $\lim_{c\to\infty}\chi(c) > 0$ Padoan, 2011)
  - Computation that is on the order of Gaussian computation
- For this distribution, $\chi(\mathbf{h})$ shows asymptotic dependence that does not approach 0 as $\mathbf{h} \to \infty$
- This occurs because all observations (near and far) share the same z_t and σ_t^2.
- We deal with this through a daily random partition (similar to Huser and Davison).

Daily random partition

- Daily random partition allows z_t and σ_t^2 to vary by site.

$$Y_t(\mathbf{s}) = X_t(\mathbf{s})\beta + \alpha z_t(\mathbf{s}) + \sigma(\mathbf{s})v_t(\mathbf{s})$$

- Consider a set of daily knots $w_{tk} \sim$ Uniform that define a random daily partition $P_{t1}, \ldots, P_{tK}$ such that

$$P_{tk} = \{s : k = \arg\min_\ell ||\mathbf{s} - w_{t\ell}||\}$$

- For $\mathbf{s} \in P_{tk}$

$$z_t(\mathbf{s}) = z_{tk}$$
$$\sigma_t^2(\mathbf{s}) = \sigma_{tk}^2$$

- Within each partition $Y_t(\mathbf{s})$ has the same MV skew-t distribution as before.

Example daily partition

Two sample partitions (number is at partition center)

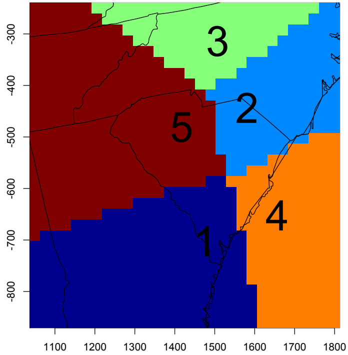

Simulated χ plots

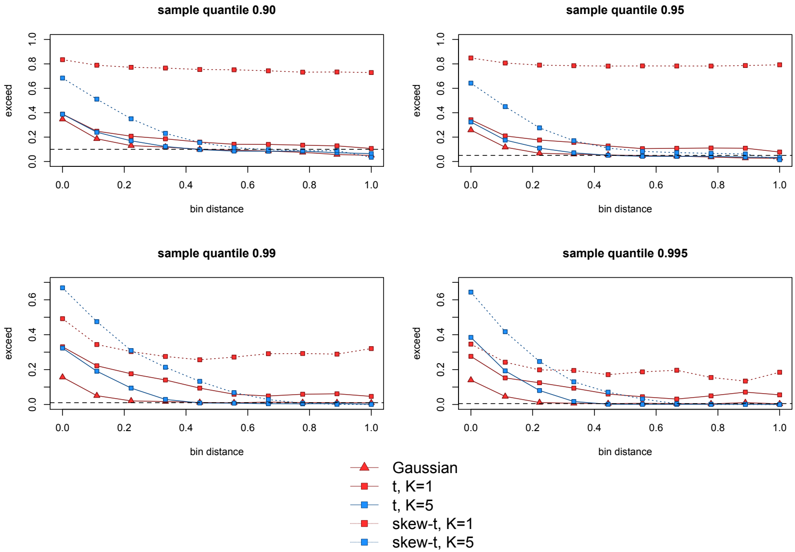

Sample simulated datasets

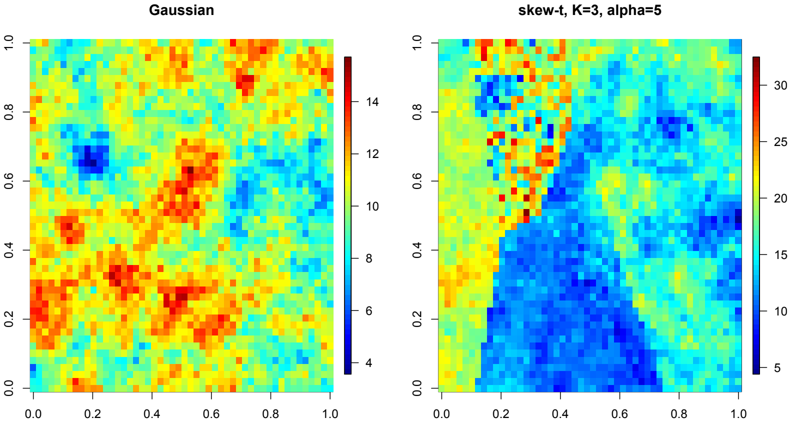

Spatiotemporal Model

- We can account for time in one of two ways
  - The mean: e.g. AR(1)
  - Three dimensional covariance model (e.g. Huser and Davison, 2014)

MCMC details

- Three main steps:
  1. Impute missing observations and censored data below T
  2. Update parameters with standard random walk Metropolis Hastings or Gibbs sampling
  3. Make spatial predictions
- Priors are selected to be conjugate when possible.

Data analysis

Ozone monitoring station locations

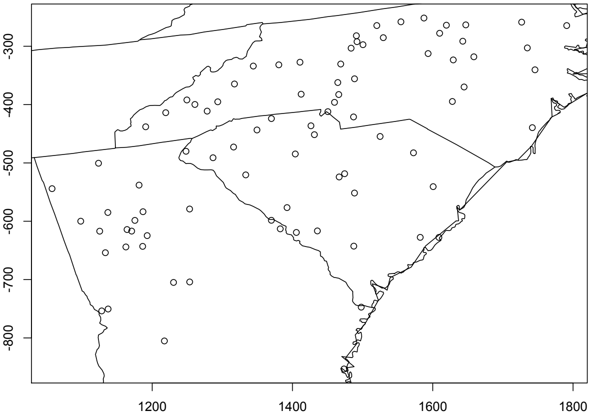

Data analysis

Max 8-hour ozone measurements at 85 sites in NC, SC, and GA for days 5 and 34.

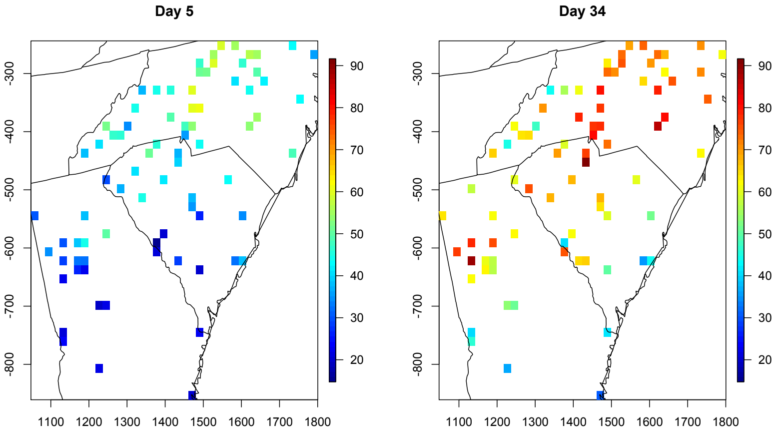

Exploratory data analysis

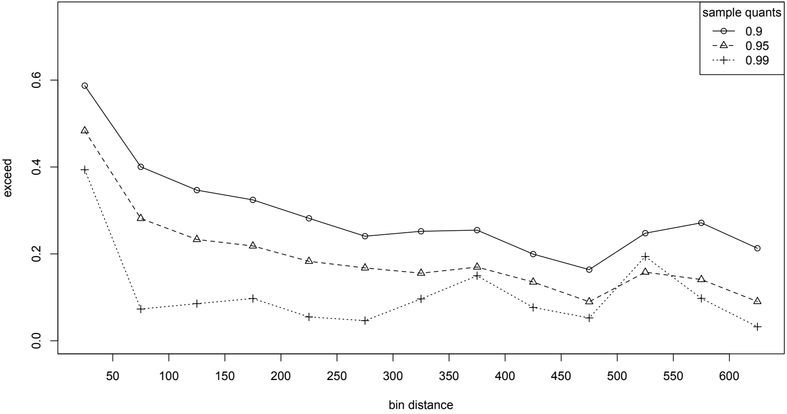

Figure: χ-plot for residuals selected ozone sample quantiles

Model comparisons

- 9 different analysis methods incorporating
  - Gaussian vs t vs skew-t marginal distribution
  - $K = 1$ partition vs $K = 5$ partitions
  - No thresholding vs thresholding at $T = 0.90$ sample quantile
- All methods use a Matérn or exponential covariance ($\nu = 0.5$)
- Compare quantile and Brier scores using 5-fold cross validation (Gneiting and Raftery, 2007)
- Mean function modeled using a first-order spatial trend

Quantile score for cross-validation

- The quantile score for the τth quantile is

$$2\{I[y < \widehat{q}(\tau)] - \tau\}(\widehat{q} - y)$$

where:

  - y is a test set value
  - $\widehat{q}(\tau)$ is the estimated τth quantile

Brier score

- The Brier score for predicting exceedance of threshold c is

$$[e(c) - P(c)]^2$$

where

  - y is a test set value
  - $e(c) = I[y > c]$
  - $P(c)$ is the predicted probability of exceeding c

Five-fold cross-validation results

			Quantile				
Marginal	K	T	0.900	0.950	0.990	0.995	0.999
Gaussian	1	0	3.807	2.981	1.797	1.493	1.044
t	1	0	2.622	1.584	0.436	0.244	0.063
t	3	0	2.643	1.601	0.455	0.265	0.085
skew-t	1	0	2.766	1.683	0.467	0.257	0.063
t	1	0.9	4.514	2.422	0.577	0.318	0.086
skew-t	1	0.9	4.396	2.388	0.573	0.311	0.079

- Brier score results are similar.

Simulation study

- 6 different data settings:
  - Gaussian vs t vs skew-t marginal distribution
  - $K = 1$ partition vs $K = 5$ partitions
- Results are similar to the results from the data analysis
- Biggest gains come from thresholding.
- Using skew models give additional gain, but small relative to gain for thresholding.

Future work

- Comparison with extreme value analysis methods
- Reporting of compliance with EPA ozone standards
- Including time in the model via standard spatiotemporal Gaussian models.

Questions

- Questions?
- Thank you for your attention.
- Acknowledgment: This work was funded by EPA STAR award R835228.

References

- Demarta, S. and McNeil, A. J. (2007) The t copula and related copulas. *International Statistical Review*, **73**, 111–129.
- Huser, R. and Davison, A. C. (2014) Space-time modelling of extreme events. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, **76**, 439–461.
- Padoan, S. A. (2011) Multivariate extreme models based on underlying skew-t and skew-normal distributions. *Journal of Multivariate Analysis*, **102**, 977–991.
- Zhang, H. and El-Shaarawi, A. (2010) On spatial skew-Gaussian processes and applications. *Environmetrics*, **21**, 33–47.